



Information Theory and Inference

# **Bayesian Optimization using Gaussian Process: Implementation from Scratch**

Eugenio Fella, Theivan Pasupathipillai, Matteo Pedrazzi,  
Gaetano Ricucci, Carlo Sgorlon Gaiatto

# Introduction

**Bayesian optimization** is a powerful tool for finding the global optimum of an expensive and black-box objective function.

It combines a probabilistic model, usually a **Gaussian process**, and an acquisition function to efficiently optimize the objective function with a small number of function evaluations.

Bayesian optimization has been successfully applied in various domains such as robotics, environmental monitoring, and automatic machine learning.

# Bayesian linear regression

Let's consider a linear regression model:

$$\mathbf{y} = \mathbf{X}\mathbf{w} \quad \text{where} \quad \mathbf{X} = \begin{pmatrix} \mathbf{x}_1 \\ \mathbf{x}_2 \\ \vdots \\ \mathbf{x}_N \end{pmatrix} \quad \text{with} \quad \mathbf{x}_i, \mathbf{w} \in \mathbb{R}^{d+1}$$

The **posterior** probability of the regression coefficients is:

$$p(\mathbf{w} \mid \mathbf{y}) = \frac{p(\mathbf{y} \mid \mathbf{w}) p(\mathbf{w})}{p(\mathbf{y})}$$

# Likelihood and prior

Let's assume that each observation is Gaussian distributed with fixed variance. Therefore the **likelihood** is:

$$p(\mathbf{y} \mid \mathbf{w}) = \mathcal{N}(\mathbf{y} \mid \mathbf{X}\mathbf{w}, \mathbf{\Sigma}) \quad \text{where} \quad \mathbf{\Sigma} = \begin{pmatrix} \sigma_1^2 & 0 & \dots & 0 \\ 0 & \sigma_2^2 & \dots & 0 \\ \vdots & \vdots & \ddots & \vdots \\ 0 & 0 & \dots & \sigma_N^2 \end{pmatrix}$$

The **prior** is chosen to be a conjugate prior, namely it is also Gaussian:

$$p(\mathbf{w}) = \mathcal{N}(\mathbf{w} \mid \mu, \mathbf{V}_0)$$

# Joint probability

The **joint probability** is the product of two multivariate Gaussian distributions that can be shown to be again Gaussian:

$$p(\mathbf{w}, \mathbf{y}) = p(\mathbf{y} \mid \mathbf{w}) p(\mathbf{w}) = \mathcal{N}(\mathbf{w}, \mathbf{y} \mid \mu', \Sigma')$$

where

$$\mu' = \begin{pmatrix} \mu \\ \mathbf{X}\mu \end{pmatrix} \quad \Sigma' = \begin{pmatrix} \mathbf{V}_0 & \mathbf{V}_0\mathbf{X}^T \\ \mathbf{X}\mathbf{V}_0 & \Sigma + \mathbf{X}\mathbf{V}_0\mathbf{X}^T \end{pmatrix}$$

# Posterior

The **posterior** can be found by starting with the joint probability and showing that the conditional probability is also Gaussian:

$$p(\mathbf{w} \mid \mathbf{y}) = \mathcal{N}(\mathbf{w} \mid \mu_{\mathbf{w}|\mathbf{y}}, \Sigma_{\mathbf{w}|\mathbf{y}})$$

where

$$\mu_{\mathbf{w}|\mathbf{y}} = \mu + \mathbf{V}_0 \mathbf{X}^T (\Sigma + \mathbf{X} \mathbf{V}_0 \mathbf{X}^T)^{-1} (\mathbf{y} - \mathbf{X} \mu)$$

$$\Sigma_{\mathbf{w}|\mathbf{y}} = \mathbf{V}_0 - \mathbf{V}_0 \mathbf{X}^T (\Sigma + \mathbf{X} \mathbf{V}_0 \mathbf{X}^T)^{-1} \mathbf{X} \mathbf{V}_0$$

# Bayesian nonparametric linear regression

This approach allows to make predictions at new locations:

$$\mathbf{y}^* = \mathbf{X}^* \mathbf{w}$$

It can be shown that the **distribution of predictions** given the observation is again Gaussian and independent of the regression coefficients:

$$p(\mathbf{y}^* \mid \mathbf{y}) = \mathcal{N}(\mathbf{y}^* \mid \mu_{\mathbf{y}^* \mid \mathbf{y}}, \Sigma_{\mathbf{y}^* \mid \mathbf{y}})$$

where

$$\mu_{\mathbf{y}^* \mid \mathbf{y}} = \mathbf{X}^* \mu + \mathbf{X}^* \mathbf{V}_0 \mathbf{X}^T (\Sigma + \mathbf{X} \mathbf{V}_0 \mathbf{X}^T)^{-1} (\mathbf{y} - \mathbf{X} \mu)$$

$$\Sigma_{\mathbf{y}^* \mid \mathbf{y}} = \mathbf{X}^* \mathbf{V}_0 \mathbf{X}^{*T} - \mathbf{X}^* \mathbf{V}_0 \mathbf{X}^T (\Sigma + \mathbf{X} \mathbf{V}_0 \mathbf{X}^T)^{-1} \mathbf{X} \mathbf{V}_0 \mathbf{X}^{*T}$$

# Kernel trick

To increase the expressiveness of the model it is common to use a non linear **feature mapping**:

$$\Phi = \phi(\mathbf{X})$$

The **kernel trick** allows us to specify an intuitive similarity between pairs of points, rather than a feature map, which in practice can be hard to define:

$$\mathbf{K}_{\mathbf{X}\mathbf{X}}(i, j) = \Phi(\mathbf{x}_i) \mathbf{V}_0 \Phi(\mathbf{x}_j)^T = \langle \Phi(\mathbf{x}_i), \Phi(\mathbf{x}_j) \rangle_{\mathbf{V}_0}$$



# Kernels

There are several stationary kernels, which are shift invariant:

$$k_{\text{MATERN}_1}(\mathbf{x}, \mathbf{x}') = \theta_0^2 \exp(-r)$$

$$k_{\text{MATERN}_3}(\mathbf{x}, \mathbf{x}') = \theta_0^2 \exp(-\sqrt{3}r)(1 + \sqrt{3}r)$$




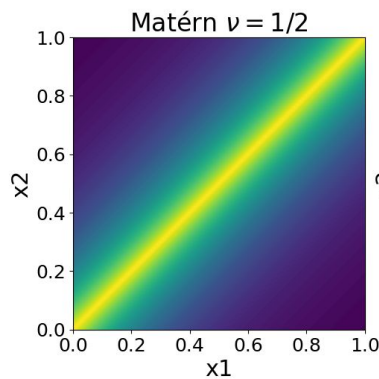

$$k_{\text{MATERN}_5}(\mathbf{x}, \mathbf{x}') = \theta_0^2 \exp(-\sqrt{5}r) \left( 1 + \sqrt{5}r + \frac{5}{3}r^2 \right)$$

$$k_{\text{Sq-exp}}(\mathbf{x}, \mathbf{x}') = \theta_0^2 \exp(-1/2r^2),$$

where

$$r^2 = (\mathbf{x} - \mathbf{x}')^T \mathbf{\Lambda} (\mathbf{x} - \mathbf{x}')$$

# Kernels



# Gaussian process

Finally we get:

$$p(\mathbf{y}^* \mid \mathbf{y}) = \mathcal{N}(\mathbf{y}^* \mid \mu_{\mathbf{y}^* \mid \mathbf{y}}, \Sigma_{\mathbf{y}^* \mid \mathbf{y}})$$

where

$$\mu_{\mathbf{y}^* \mid \mathbf{y}} = \Phi^* \mu + \mathbf{K}_{\mathbf{X}^* \mathbf{X}} \hat{\mathbf{K}}_{\mathbf{X} \mathbf{X}}^{-1} (\mathbf{y} - \Phi \mu)$$

$$\Sigma_{\mathbf{y}^* \mid \mathbf{y}} = \mathbf{K}_{\mathbf{X}^* \mathbf{X}^*} - \mathbf{K}_{\mathbf{X}^* \mathbf{X}} \hat{\mathbf{K}}_{\mathbf{X} \mathbf{X}}^{-1} \mathbf{K}_{\mathbf{X} \mathbf{X}^*}$$

# Acquisition functions

There are many selection strategies that utilize the posterior model to select the next query point:

**Probability of Improvement (PI):** it aims to maximize the probability of finding a point that is better than the current best solution.

$$\text{PI}(x) = P(f(x) \geq f(x_{\text{best}}))$$

# Acquisition functions

**Expected Improvement (EI):** it is defined as the expected value of the improvement over the current best solution.

$$\text{EI}(x) = \mathbb{E} [\max(f(x) - f(x_{\text{best}}), 0)]$$

**Upper Confidence Bound (UCB):** it is defined as the sum of the mean function value and a measure of the uncertainty in the function value.

$$\text{UCB}(x) = \mu(x) + \kappa\sigma(x)$$

# Hyperparameter tuning: Maximum Likelihood Estimation

**Maximizing the likelihood** of the kernel parameters given the data can give an estimate on the best parameters to choose for the kernel:

$$\hat{\theta} = \arg \max_{\theta} \log p(\mathbf{y} \mid \mathbf{X}, \theta)$$

The typical estimation of the hyperparameters by maximizing the marginal likelihood can easily fall into traps. The optimization problem was tackled by **minimizing the negative log-likelihood**:

$$\log p(\mathbf{y} \mid \mathbf{X}, \theta) = -\frac{1}{2} \log |\mathbf{K}(\theta) + \sigma^2 \mathbf{I}| - \frac{1}{2} \mathbf{y}^T (\mathbf{K}(\theta) + \sigma^2 \mathbf{I})^{-1} \mathbf{y} - \frac{n}{2} \log 2\pi$$

# Hyperparameter tuning: gradient descent

Minimize the negative log-likelihood using **gradient descent** is another approach to obtain an estimate of the best hyperparameters to use:

$$\theta_l^{(t+1)} = \theta_l^{(t)} - \lambda \nabla \ell(\boldsymbol{\theta}; \mathbf{X}_n, \mathbf{y}_n)_l$$

Each gradient component of the negative log-likelihood can be computed with the following formula:

$$\nabla \ell(\boldsymbol{\theta}; \mathbf{X}_n, \mathbf{y}_n)_l = \frac{1}{2n} \left[ -\mathbf{y}_n^\top \mathbf{K}_n^{-1} \frac{\partial \mathbf{K}_n}{\partial \theta_l} \mathbf{K}_n^{-1} \mathbf{y}_n + \text{tr} \left( \mathbf{K}_n^{-1} \frac{\partial \mathbf{K}_n}{\partial \theta_l} \right) \right]$$

# Hyperparameter tuning: marginal acquisition function

Point estimation methods like MLE provide a single estimate but do not account for the uncertainty in the estimate.

**Marginalization** provides a distribution over the hyperparameters that consider this uncertainty:

$$\alpha(x) = \mathbb{E}_{\theta | \mathcal{D}_n} [\alpha(x; \theta)] \approx \frac{1}{M} \sum_{i=1}^M \alpha \left( x; \theta_n^{(i)} \right)$$

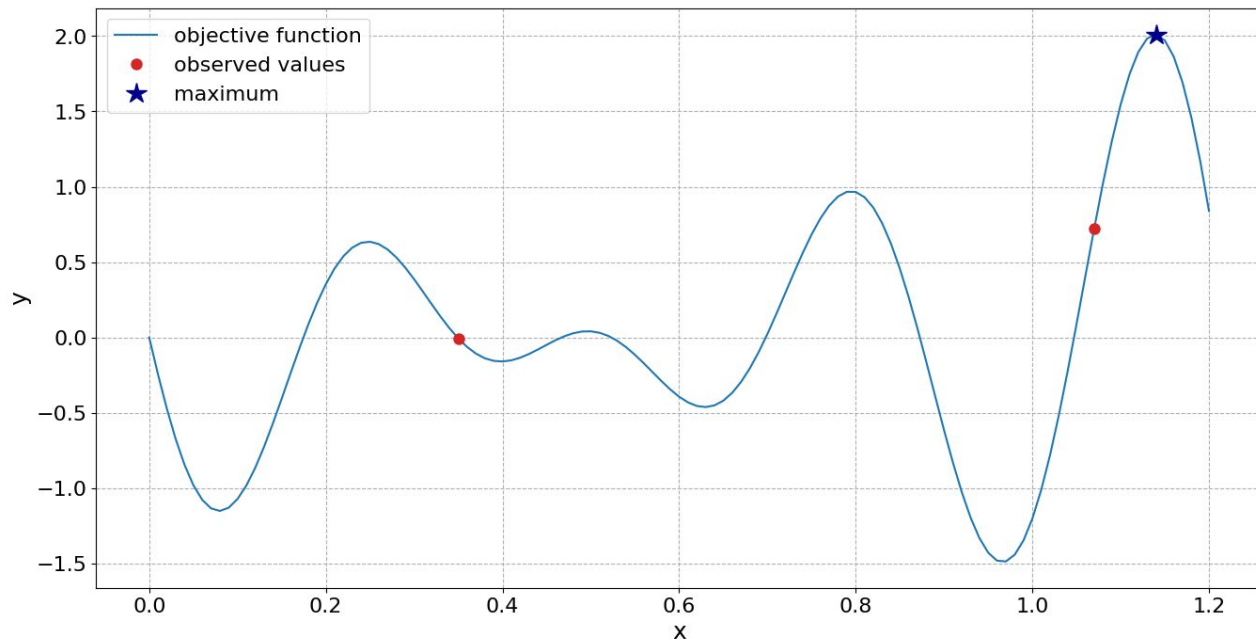
The sampling procedure is not trivial given that it requires to tune the Metropolis-Hasting hyperparameters.



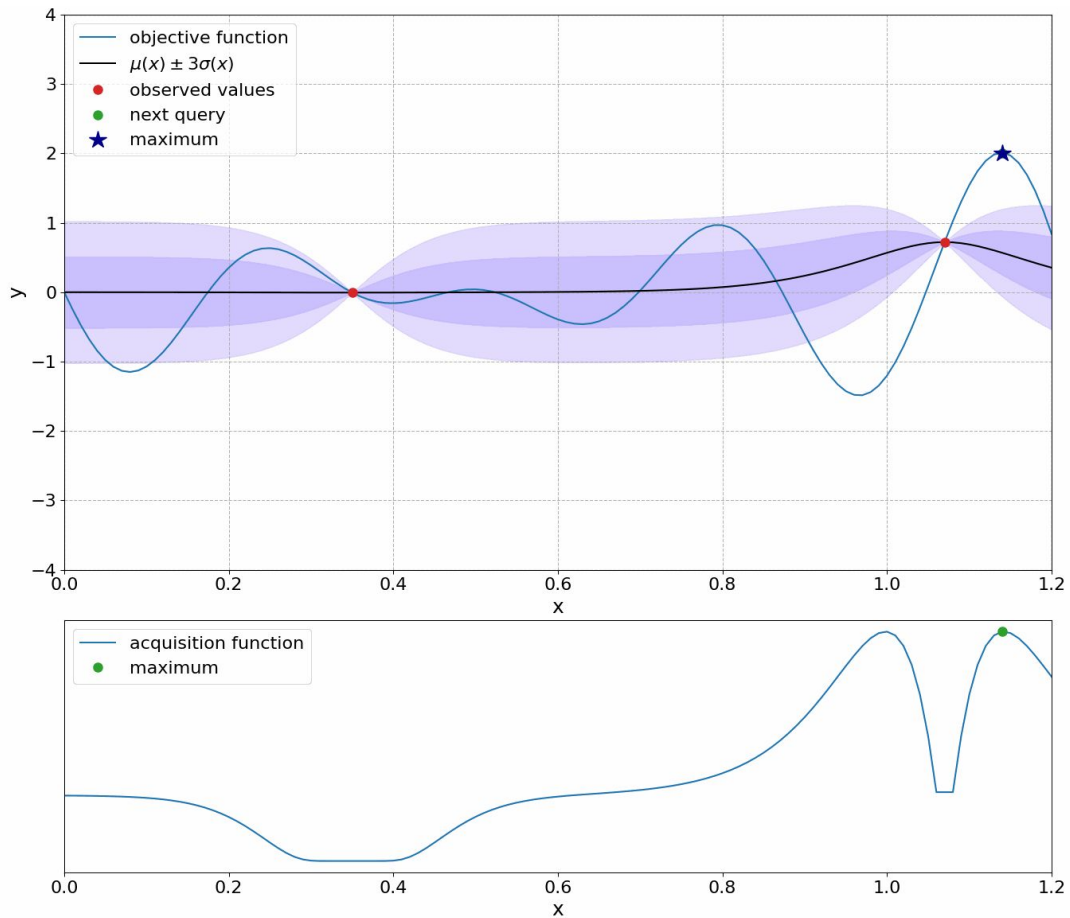
# 1D analytic function

$$y = -(a - bx) \sin(cx)$$

$$[a = 1.3, b = 3, c = 18]$$




# 1D analytic function: simulation




# 1D analytic function: hyperparameter tuning

Comparison of hyperparameter tuning methods:




# 1D analytic function: acquisition functions

Comparison of acquisition functions:




# 2D analytic function : Branin-Hoo

$$f(x_1, x_2) = -a(x_2 - bx_1^2 + cx_1 - r)^2 - s(1 - t)\cos(x_1) - s$$
$$[a = 1, b = 5.1/(4\pi^2), c = 5/\pi, r = 6, s = 10, t = 1/(8\pi)]$$




# 2D analytic function: simulation




# 2D analytic function: hyperparameter tuning

Comparison of hyperparameter tuning methods:



# 2D analytic function: acquisition functions

Comparison of acquisition functions:





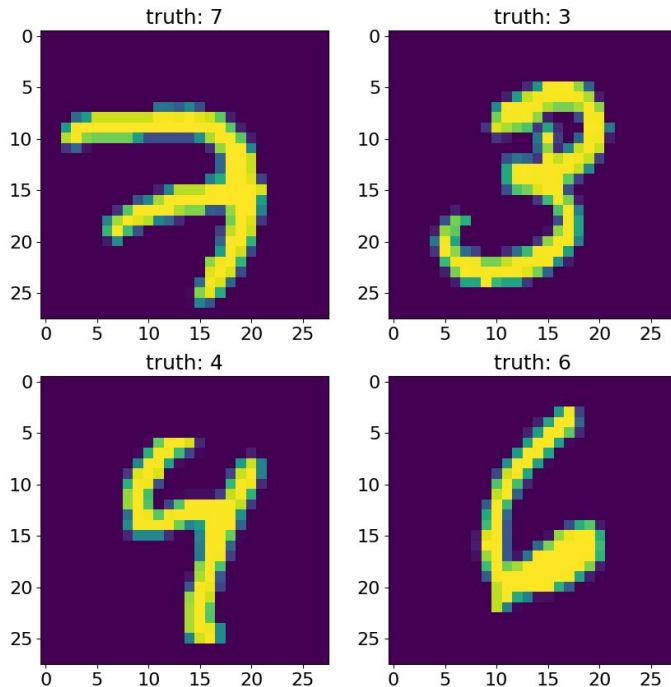
# Multilayer perceptron

We trained a simple fully connected feedforward **neural network** (NN) on the *mnist-784* dataset from OpenML.

The number of neurons in each hidden layer is set to 5 and we look for the maximum scores obtained by the NN, exploring two hyperparameters:

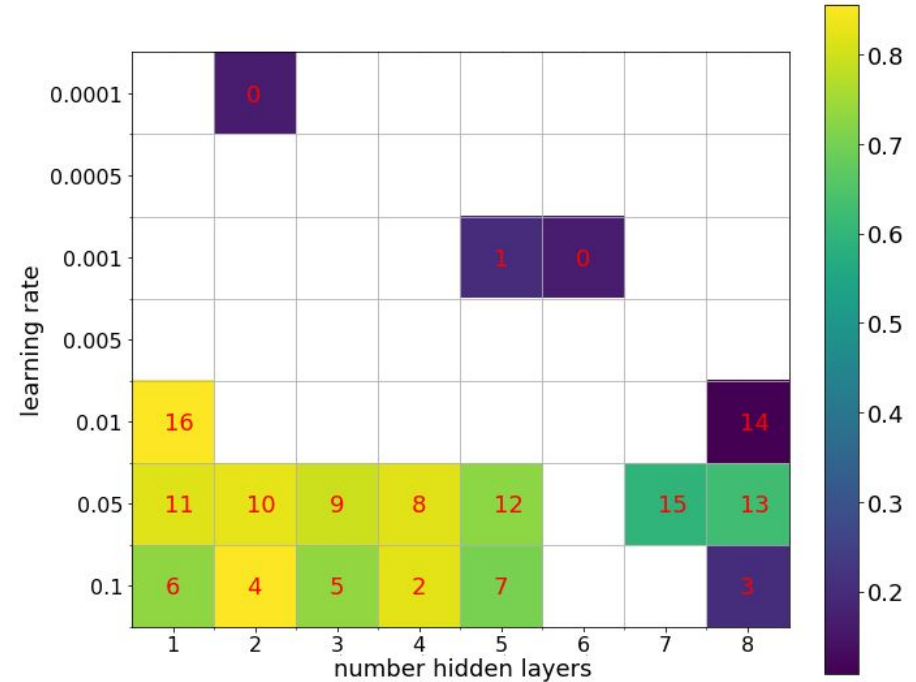
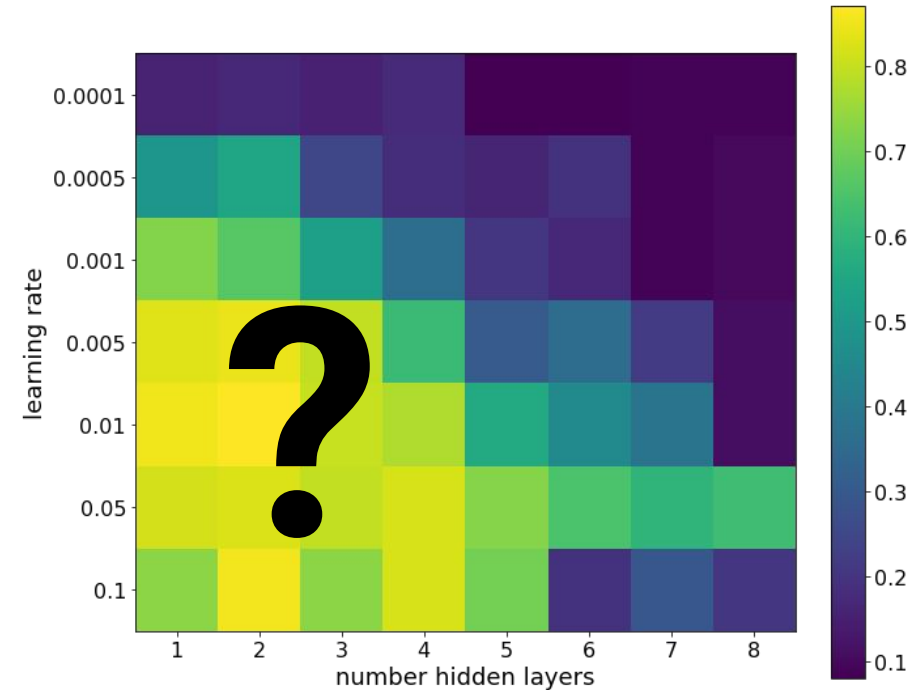
- number of hidden layers
- learning rate

*mnist-784* dataset



# Multilayer perceptron: simulation

Bayesian optimization



# Multilayer perceptron: hyperparameter tuning

Comparison of hyperparameter tuning methods:

